

AI 221: Machine Exercise 4

Instructions:

- Read and answer each problem using computer code. This MEX should be done *individually*.
- Each item should be answered as either a Jupyter Notebook or a MATLAB Live Script, exported as a single PDF file for the entire MEX. Make sure to HIGHLIGHT your final answers.
- When done, submit the PDF file through UVLE.

8x8 Handwritten Digits Visualization and Classification

For this problem, we'll use the example given in Lecture 4: 8x8 handwritten digits data. Load the `load_digits` data from `sklearn.datasets`.

Do the following:

- a. [30 pts] Normalize the X data using Standard Scaler. Then, project all the X data into 2 dimensions using 6 dimensionality reduction techniques:
 - i. Local Linear Embedding (`n_neighbors = 200, random_state = 0`)
 - ii. t-SNE (`perplexity = 50, random_state = 0`)
 - iii. Isomap (`n_neighbors = 200`)
 - iv. Laplacian Eigenmap (`n_neighbors = 200`)
 - v. Kernel PCA (`kernel = 'rbf', gamma = 0.01`)
 - vi. PCA

The points should then be colored according to the digit labels, `y`.

Which of the methods produced clear clusters of data points?

- b. [10 pts] For both KPCA and PCA in the previous item, report the CPV plot and determine the number of principal components needed to retain 95% cumulative percent variance.
- c. In this item, we will perform classification with and without dimensionality reduction. First, split the data into 70% training and 30% testing, *stratified* according to class label. Compare the results between the following methods:
 - i. [20 pts] Make a pipeline using StandardScaler, Kernel PCA (`kernel = 'sigmoid', n_components=40`), and SVC (default hyper-parameters). Fit the pipeline on the training set, then report the accuracy and F1-score on the test set.
 - ii. [20 pts] Make a pipeline using StandardScaler, Linear Discriminant Analysis (`n_components=9`), and SVC (default hyper-parameters). Fit the pipeline on the training set, then report the accuracy and F1-score on the test set. Can we extract more than 9 components from LDA? Why or why not?
 - iii. [20 pts] Make a pipeline using StandardScaler and SVC (default hyper-parameters) alone. Fit the pipeline on the training set, then report the accuracy and F1-score on the test set. This is the case without any dimensionality reduction.

Which of the three methods had the best test performance for classification? Explain.

END OF EXERCISE