

MP5

December 2, 2025

1 Machine Exercise 5

Jhon Mark L. Acopiado

University of the Philippines Diliman

Masters of Engineering in Artificial Intelligence

AI221:Classical Machine Learning

1.1 Part 1: Country Data from HELP International

1.1.1 1.I. Load Dataset

```
[1]: # Source Code From : https://www.kaggle.com/datasets/rohan0301/
      ↵unsupervised-learning-on-country-data/data
# Install dependencies as needed:
# pip install kagglehub[pandas-datasets]
import kagglehub
from kagglehub import KaggleDatasetAdapter

# Set the path to the file you'd like to load
file_path = "Country-data.csv"

# Load the latest version
df = kagglehub.load_dataset(
    KaggleDatasetAdapter.PANDAS,
    "rohan0301/unsupervised-learning-on-country-data",
    file_path,
    # Provide any additional arguments like
    # sql_query or pandas_kwargs. See the
    # documentation for more information:
    # https://github.com/Kaggle/kagglehub/blob/main/README.
    ↵md#kaggledatasetadapterpandas
)

print("First 5 records:", df.head())
```

```
C:\Users\jhon\AppData\Local\Temp\ipykernel_20780\2839170658.py:11:
DeprecationWarning: Use dataset_load() instead of load_dataset(). load_dataset()
```

```

will be removed in a future version.

df = kagglehub.load_dataset()

First 5 records:
   country child_mort exports health imports
income \
0    Afghanistan     90.2     10.0     7.58    44.9    1610
1      Albania       16.6     28.0     6.55    48.6    9930
2    Algeria        27.3     38.4     4.17    31.4   12900
3      Angola       119.0     62.3     2.85    42.9    5900
4 Antigua and Barbuda     10.3     45.5     6.03    58.9   19100

   inflation life_expec total_fer gdpp
0      9.44      56.2      5.82    553
1      4.49      76.3      1.65   4090
2     16.10      76.5      2.89   4460
3     22.40      60.1      6.16   3530
4      1.44      76.8      2.13  12200

```

```
[2]: # Prompt: Store Features and Target Variable
X = df.drop(columns=['country'])
Y = df['country']
```

1.1.2 1.II. K-Means

Item 1.a. Normalize the features data using Standard Scaler. Then, perform K-means clustering on all features. Display the elbow plot (Inertia vs. no. of clusters) and the silhouette score plot. What number of clusters is recommended?

```
[3]: # Prompt: Normalize the features of X using StandardScaler
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
import numpy as np
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

import matplotlib.pyplot as plt

# Define the range for the number of clusters
cluster_range = range(2, 11)
inertia = []
silhouette_scores = []

# Perform KMeans clustering for each number of clusters
for n_clusters in cluster_range:
    kmeans = KMeans(n_clusters=n_clusters, random_state=42)
    kmeans.fit(X_scaled)
    inertia.append(kmeans.inertia_)
```

```

silhouette_scores.append(silhouette_score(X_scaled, kmeans.labels_))

def plot_elbow_and_silhouette(cluster_range, inertia, silhouette_scores):
    # Plot the elbow plot (Inertia vs. number of clusters)
    plt.figure(figsize=(12, 5))

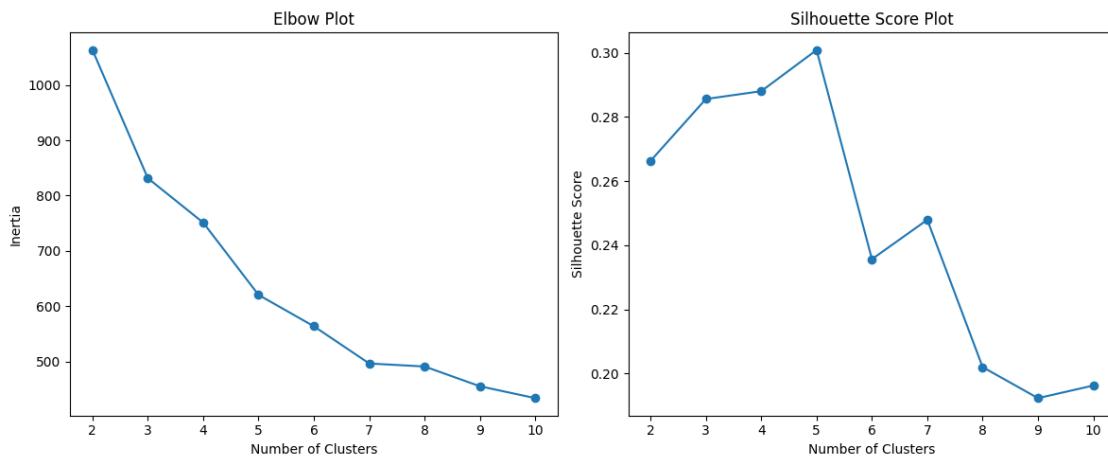
    plt.subplot(1, 2, 1)
    plt.plot(cluster_range, inertia, marker='o')
    plt.title('Elbow Plot')
    plt.xlabel('Number of Clusters')
    plt.ylabel('Inertia')

    # Plot the silhouette score plot
    plt.subplot(1, 2, 2)
    plt.plot(cluster_range, silhouette_scores, marker='o')
    plt.title('Silhouette Score Plot')
    plt.xlabel('Number of Clusters')
    plt.ylabel('Silhouette Score')

    plt.tight_layout()
    plt.show()

# Call the function
plot_elbow_and_silhouette(cluster_range, inertia, silhouette_scores)

```



Item 1.a. What number of clusters is recommended?

Based on the elbow plot and Silhouette Score, the recommended number of clusters is either 5 or 7. For both 5 and 7 clusters, the Silhouette Score increases, indicating that the data points are well-clustered and distinct from other clusters. The elbow plot shows the relationship between the number of clusters and the inertia. At 7 clusters, the rate of decrease in inertia slows down

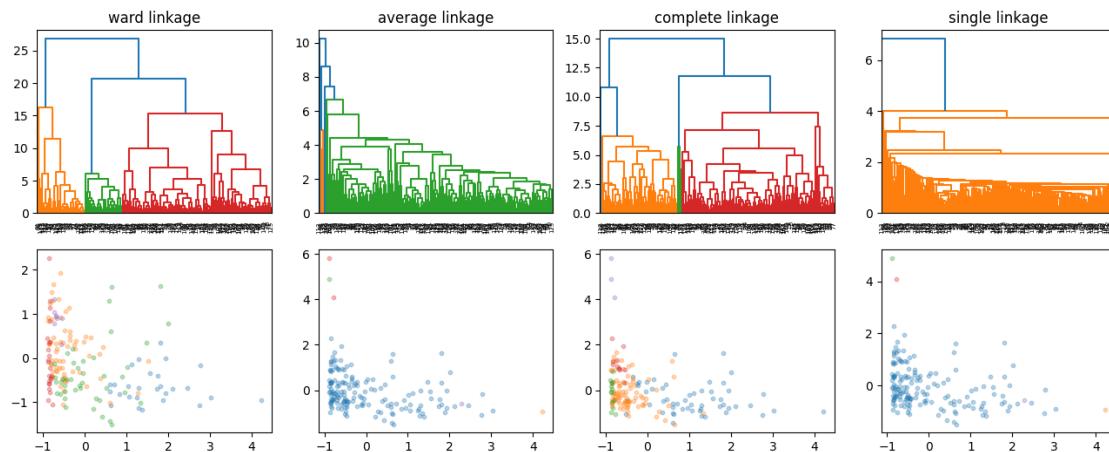
significantly compared to 5 clusters.

If I were to choose only one, I would select 7 clusters.

1.1.3 1.III. Hierarchical clustering

Item 1.b. Perform hierarchical clustering on the normalized data set and compare the results of various linkage methods. Which one would you recommend? Why is this recommended clustering informative?

```
[4]: # Soure Code is From AI221 Github Repository:  
    ↪AI221\Clustering_Anomaly_Detect\kmeans_aggro_blobs.ipynb  
from scipy.cluster.hierarchy import linkage, dendrogram, cut_tree  
from time import time  
  
fig3 = plt.figure(figsize=(16,6))  
ctr = 1  
t0 = time()  
for method in ('ward', 'average', 'complete', 'single'):  
    Z = linkage(X_scaled, method=method)  
    ax = fig3.add_subplot(240 + ctr)  
    ax.set_title("%s linkage" % method)  
    dendrogram(Z)  
    cutree = cut_tree(Z, n_clusters=7).flatten()  
    ax = fig3.add_subplot(240 + ctr + 4)  
    for j in range(0,5):  
        ax.scatter(X_scaled[cutree == j,0], X_scaled[cutree == j,1], s=10,  
    ↪alpha=0.3)  
    ctr += 1  
  
plt.show()  
print(f"Elapsed Time: {time()-t0} sec")
```



Elapsed Time: 3.20088267326355 sec

Item 1.b. Which one would you recommend? Why is this recommended clustering informative?

I recommend using the Ward linkage method. For this analysis, I set cut_tree to 7 clusters to align with my response to the previous item. The visualization demonstrates that the Ward linkage method is more informative compared to other methods. Ward linkage minimizes the variance within clusters, resulting in compact and well-separated clusters. This makes it particularly effective when the data has a Gaussian-like distribution or when compact clusters are desired. The method ensures that the clusters formed are homogeneous and interpretable, making it a suitable choice for this dataset.

1.1.4 1.IV. PCA (2 features) + K-means

Item 1.c. This time, perform PCA to reduce the features data into 2D. Then, perform K-means clustering on the reduced data set just as in item (a). Display the elbow plot and silhouette score plot as well.

```
[5]: # Prompt: Perform PCA on X to reduce its dimensions to 2 components
from sklearn.decomposition import PCA
import random
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X)

# Prompt: Display a scatter plot of the PCA-reduced data with some points labeled
plt.figure(figsize=(8, 6))
plt.scatter(X_pca[:, 0], X_pca[:, 1], alpha=0.7)
# Prompt: Randomly select some countries to label
selected_countries = random.sample(list(df['country']), 20)

for i, country in enumerate(df['country']):
    if country in selected_countries:
        plt.text(X_pca[i, 0], X_pca[i, 1], s=country, fontsize=8)
plt.title('PCA of Country Data')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()
# Repeat KMeans clustering on PCA-reduced data
cluster_range = range(2, 11)
inertia = []
silhouette_scores = []

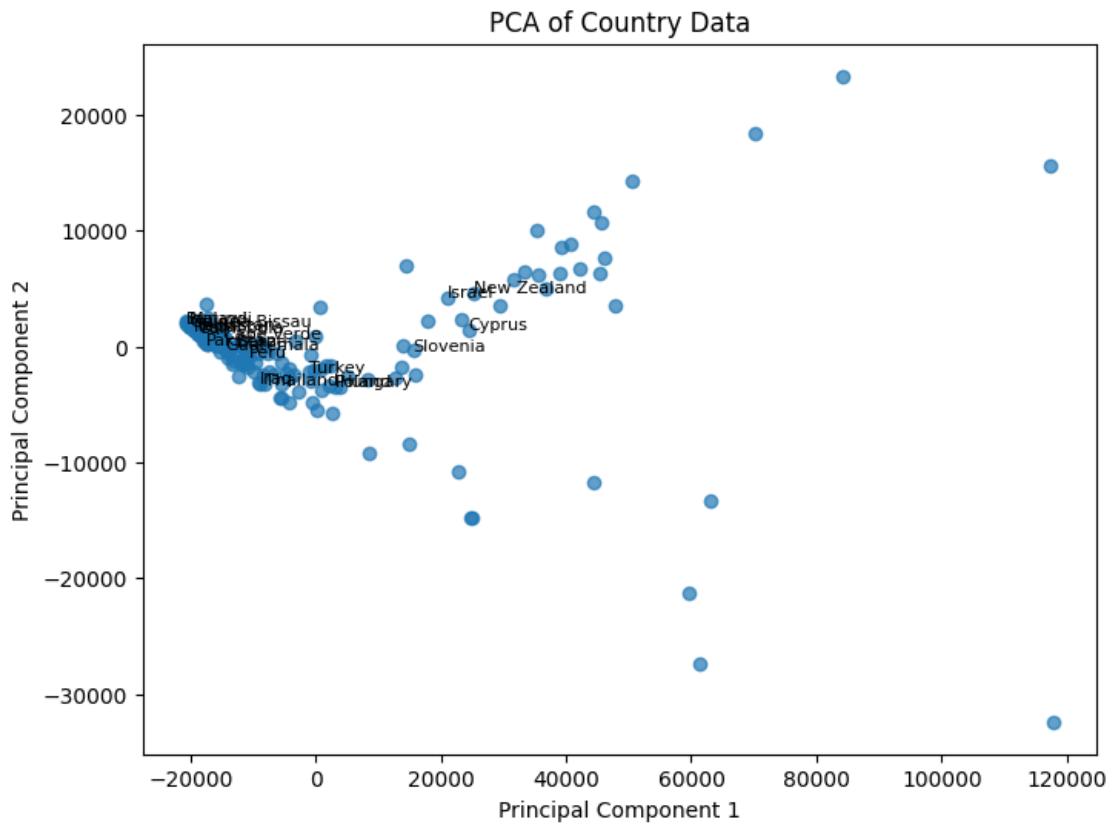
# Perform KMeans clustering for each number of clusters
for n_clusters in cluster_range:
```

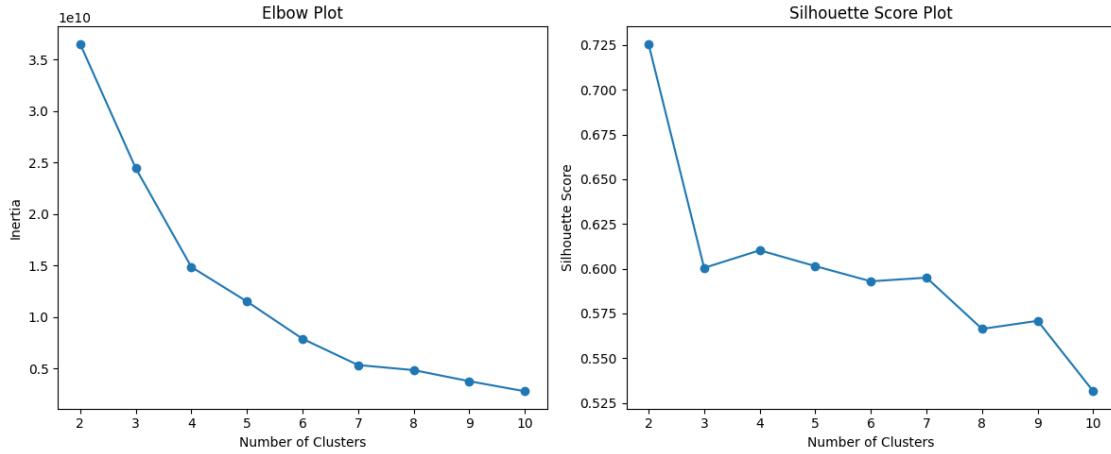
```

kmeans = KMeans(n_clusters=n_clusters, random_state=42)
kmeans.fit(X_pca)
inertia.append(kmeans.inertia_)
silhouette_scores.append(silhouette_score(X_pca, kmeans.labels_))

plot_elbow_and_silhouette(cluster_range, inertia, silhouette_scores)

```





Item 1.c. What number of clusters is recommended? You can label some countries in the 2D mapping.

I would recommend 7 clusters as it is the number where the elbow plot shows a significant decrease in inertia. Additionally, the Silhouette score for 7 clusters is higher than that of its neighboring cluster numbers, indicating well-defined and distinct clusters.

Item 1.d. Based on the recommended no. of clusters in item (b), make your own descriptions of each cluster. What range of values of features are unique to each cluster?

To support my answer, let me get insights from the table, and plots below. I will base my descriptions on the mean values of the feature variables for each cluster. Below are the descriptions for each cluster:

- **Cluster 1:** This cluster is characterized by high health values (10.51) and low inflation (1.33). These countries likely have strong healthcare systems and stable economies.
- **Cluster 2:** Defined by the lowest income (3017.69), life expectancy (59.04), and GDP per capita (1455.76). These countries may face significant economic and health challenges.
- **Cluster 3:** Notable for the highest child mortality (130.00), inflation (104.00), and fertility rates (5.84). These countries may have developing economies with high population growth.
- **Cluster 4:** Distinguished by the highest exports (176.00), imports (156.67), life expectancy (81.43), and GDP per capita (57566.67). These countries are likely developed nations with strong trade and high standards of living.
- **Cluster 5:** Characterized by the highest income (67171.43) and lowest health values (3.28). These countries may have high wealth but face challenges in healthcare access or outcomes.
- **Cluster 0:** Represents countries with average values across most features, indicating balanced development without extremes in any specific area.
- **Cluster 6:** Defined by unique outliers or anomalies in the dataset, possibly representing countries with rare or exceptional conditions not captured by other clusters.

The following blocks of code support the analysis by providing detailed insights into the clustering results:

- Mean Values of Features for Each Cluster: The code calculates the mean values of each feature for every cluster, allowing us to identify patterns and differences across clusters.

- Box Plots for Feature Distributions: The code generates box plots to visualize the distribution of each feature within the clusters. This helps in understanding the variability and spread of the data for each feature across different clusters.
- Clusters with Highest and Lowest Mean Values: The code identifies which cluster has the highest and lowest mean value for each feature, highlighting the distinguishing characteristics of the clusters.
- Visualization of Cluster Grouping: Finally, the code includes a plot that visually represents how the clusters group as the number of clusters increases. This provides a clear picture of the clustering structure and how the data is segmented.

```
[6]: import seaborn as sns

cluster_labels = KMeans(n_clusters=7, random_state=42).fit_predict(X_scaled)
df_with_clusters = df.copy()
df_with_clusters['Cluster'] = cluster_labels # Add cluster labels to the DataFrame
cluster_summary = df_with_clusters.drop(columns=['country']).groupby('Cluster').agg(['mean'])
print(display(cluster_summary))

# Prompt: Plot all feature distributions by cluster in the same figure
plt.figure(figsize=(20, 15))
num_features = len(df_with_clusters.columns[1:-1]) # Exclude 'country' and 'Cluster' columns
for i, feature in enumerate(df_with_clusters.columns[1:-1], 1):
    plt.subplot((num_features + 2) // 3, 3, i) # Arrange subplots in a grid
    for cluster in sorted(df_with_clusters['Cluster'].unique()):
        cluster_data = df_with_clusters[df_with_clusters['Cluster'] == cluster]
        plt.boxplot(cluster_data[feature], positions=[cluster], widths=0.6, patch_artist=True)
        plt.title(f'{feature.capitalize()} Distribution by Cluster')
        plt.xlabel('Cluster')
        plt.ylabel(feature.capitalize())

plt.tight_layout()
plt.show()

# Prompt: Determine which cluster has the highest and lowest values for each column
highest_lowest_clusters = {}

for column in df_with_clusters.columns[1:-1]: # Exclude 'country' and 'Cluster' columns
    cluster_means = df_with_clusters.groupby('Cluster')[column].mean()
    highest_cluster = cluster_means.idxmax()
    lowest_cluster = cluster_means.idxmin()
```

```

highest_lowest_clusters[column] = {
    'highest_cluster': highest_cluster,
    'highest_value': cluster_means[highest_cluster],
    'lowest_cluster': lowest_cluster,
    'lowest_value': cluster_means[lowest_cluster]
}

# Display the results
for feature, clusters in highest_lowest_clusters.items():
    print(f"Feature: {feature}")
    print(f" Highest Cluster: {clusters['highest_cluster']} (Value:{clusters['highest_value']:.2f})")
    print(f" Lowest Cluster: {clusters['lowest_cluster']} (Value:{clusters['lowest_value']:.2f})")

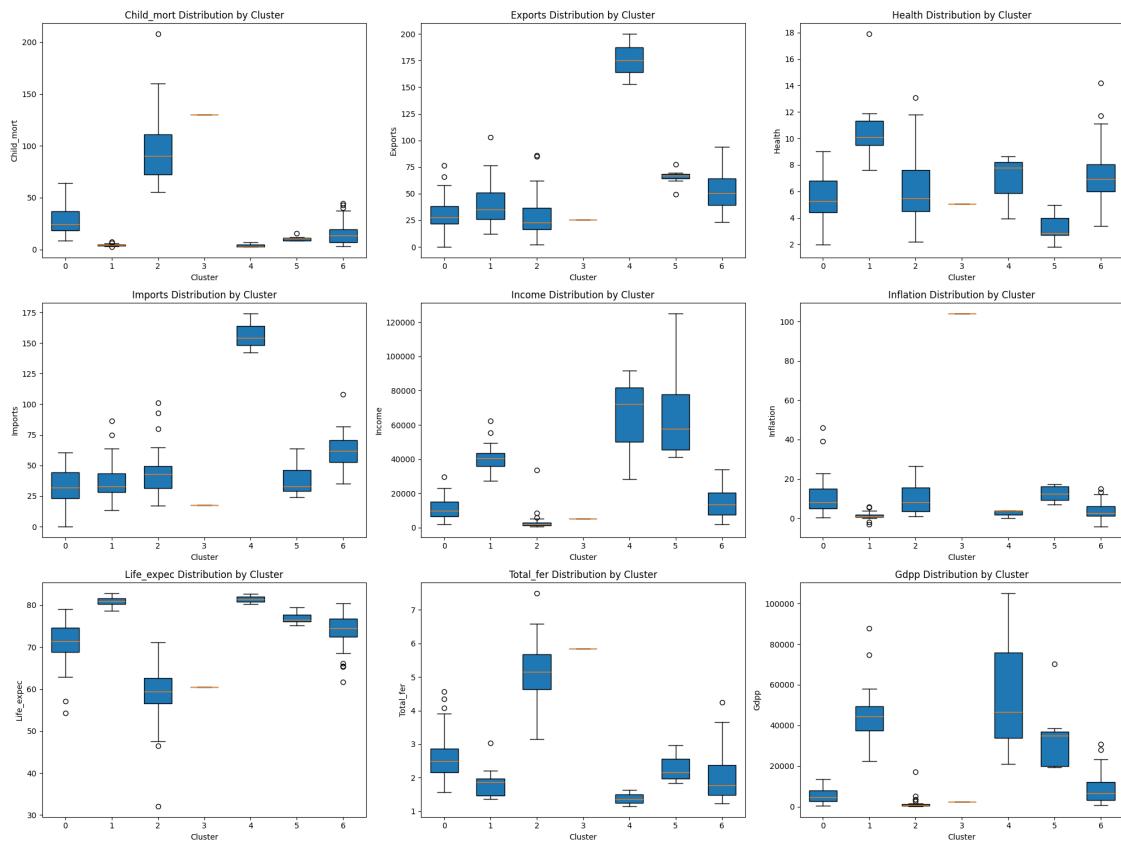
# Prompt: Visualize KMeans clustering results for cluster counts from 2 to 10
inertia_values = []
silhouette_avgs = []
fig2 = plt.figure(figsize=(20,20))
for i in range(2,11):
    kmeans = KMeans(n_clusters=i, n_init=10).fit(X_scaled)
    cluster_labels = kmeans.labels_
    centroids = kmeans.cluster_centers_
    inertia_values.append(kmeans.inertia_)
    silhouette_avg = silhouette_score(X_scaled, cluster_labels)
    silhouette_avgs.append(silhouette_avg)
    ax = fig2.add_subplot(330 + i - 1)
    for j in range(0,i):
        ax.scatter(X_scaled[cluster_labels == j,0], X_scaled[cluster_labels == j,1], s=10, alpha=0.3)
    ax.scatter(centroids[:,0],centroids[:,1], s=50, color='b', edgecolor='k')
    ax.set_title(f"Inertia = {kmeans.inertia_}\nSilhouette Score = {silhouette_avg}\nNumber of clusters = {i}")

```

	child_mort mean	exports mean	health mean	imports mean	income mean	\
Cluster						
0	29.430435	30.621500	5.569565	33.577520	10946.956522	
1	4.295652	40.730435	10.513478	38.247826	40265.217391	
2	96.009524	27.988333	6.334286	43.811905	3017.690476	
3	130.000000	25.300000	5.070000	17.400000	5150.000000	
4	4.133333	176.000000	6.793333	156.666667	64033.333333	
5	10.700000	65.557143	3.281429	38.700000	67171.428571	
6	15.306667	51.824444	7.238889	62.400000	14206.888889	
	inflation life_expec total_fer			gdpp		

Cluster	mean	mean	mean	mean
0	10.601435	71.139130	2.590652	5526.000000
1	1.334913	80.891304	1.810870	45417.391304
2	9.820357	59.042857	5.129286	1455.761905
3	104.000000	60.500000	5.840000	2330.000000
4	2.468000	81.433333	1.380000	57566.666667
5	12.517143	76.928571	2.287143	34057.142857
6	3.771511	73.928889	2.001556	8703.244444

None



Feature: child_mort

Highest Cluster: 3 (Value: 130.00)

Lowest Cluster: 4 (Value: 4.13)

Feature: exports

Highest Cluster: 4 (Value: 176.00)

Lowest Cluster: 3 (Value: 25.30)

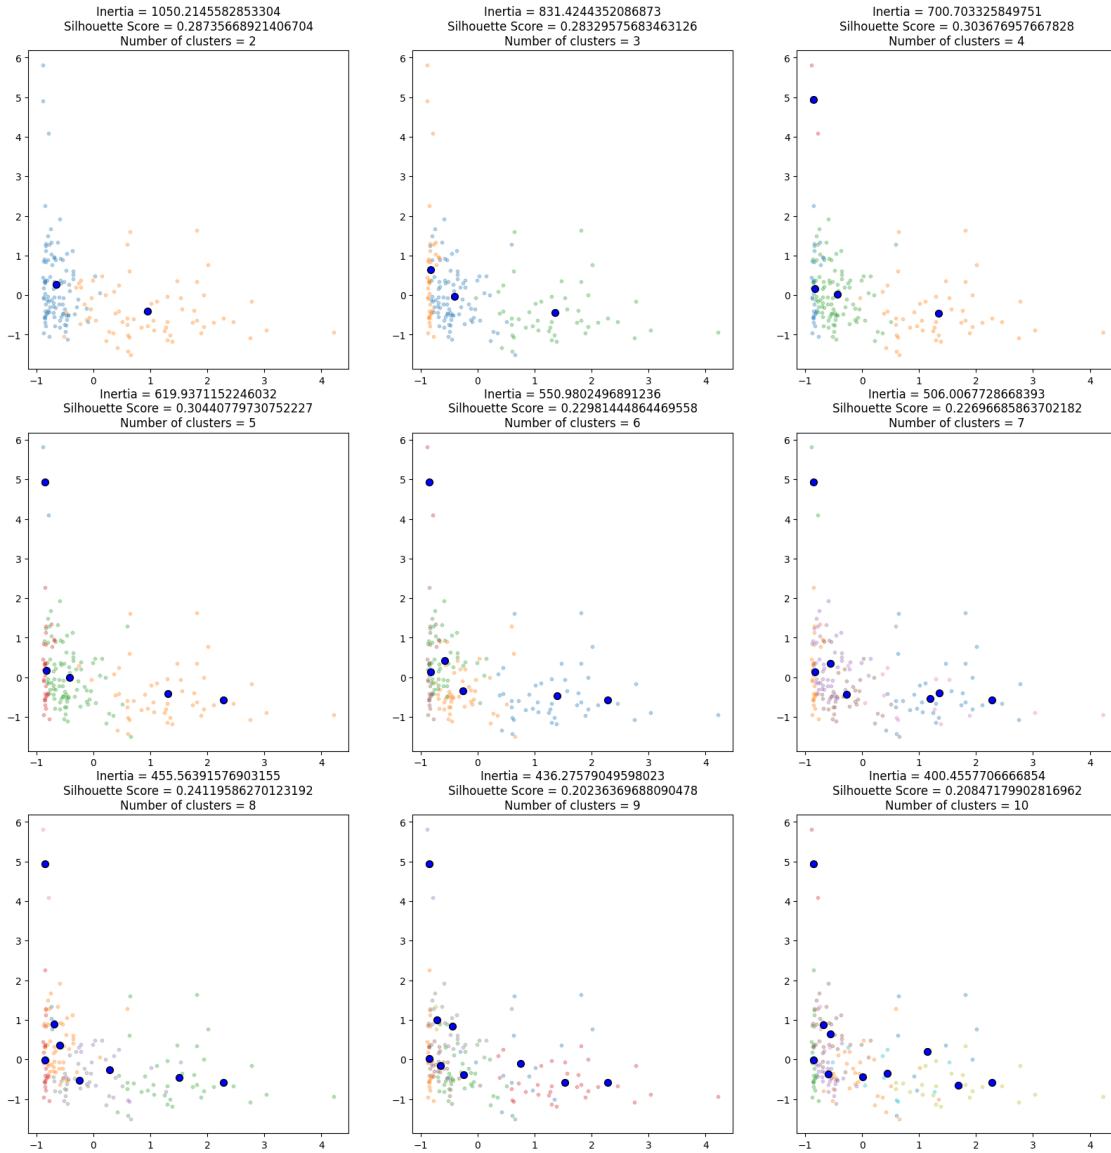
Feature: health

Highest Cluster: 1 (Value: 10.51)

Lowest Cluster: 5 (Value: 3.28)

Feature: imports

```
Highest Cluster: 4 (Value: 156.67)
Lowest Cluster: 3 (Value: 17.40)
Feature: income
    Highest Cluster: 5 (Value: 67171.43)
    Lowest Cluster: 2 (Value: 3017.69)
Feature: inflation
    Highest Cluster: 3 (Value: 104.00)
    Lowest Cluster: 1 (Value: 1.33)
Feature: life_expec
    Highest Cluster: 4 (Value: 81.43)
    Lowest Cluster: 2 (Value: 59.04)
Feature: total_fer
    Highest Cluster: 3 (Value: 5.84)
    Lowest Cluster: 4 (Value: 1.38)
Feature: gdpp
    Highest Cluster: 4 (Value: 57566.67)
    Lowest Cluster: 2 (Value: 1455.76)
```



1.1.5 1.V. 2D PCA + Anomaly Detection

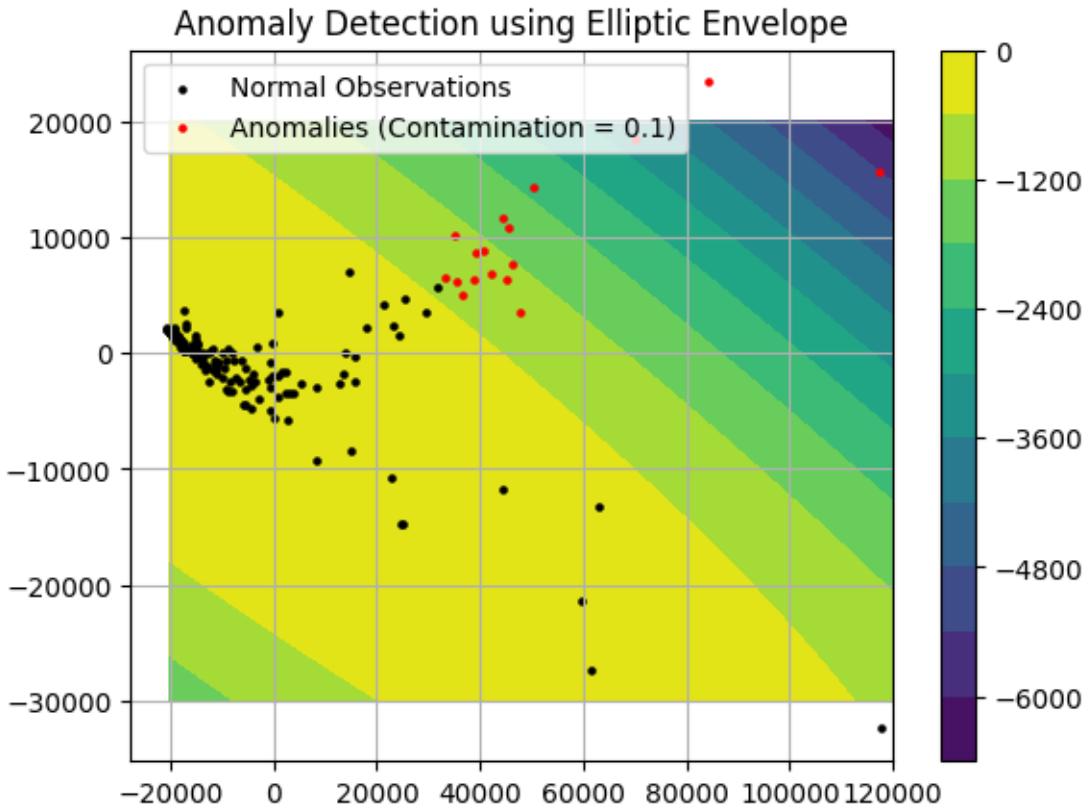
Item 1.e. Based on the 2D PCA mapping, perform anomaly detection using any method.

```
[8]: import matplotlib.pyplot as plt
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.covariance import EllipticEnvelope
from sklearn.neighbors import LocalOutlierFactor
from sklearn.neighbors import KernelDensity
from sklearn.ensemble import IsolationForest
from sklearn.datasets import make_blobs
```

```
from sklearn.svm import OneClassSVM
```

1.V.a. Elliptic Envelope

```
[9]: # Source Code is From AI221 Github Repository:  
    ↪AI221\Clustering_Anomaly_Detect\anomaly_detect_methods.ipynb  
Xp, Yp = np.meshgrid(np.linspace(-20000,120000),np.linspace(-30000,20000))  
XY = np.vstack([Xp.ravel(), Yp.ravel()]).T  
  
envelope = EllipticEnvelope(random_state=0, contamination=0.1).fit(X_pca)  
Zp = envelope.score_samples(XY)  
Zp = Zp.reshape(Xp.shape)  
  
# Get the anomalous data points  
y_pred = envelope.predict(X_pca)  
normals = X_pca[y_pred == 1,:]  
anomals = X_pca[y_pred == -1,:]  
  
cntr = plt.contourf(Xp, Yp, Zp, levels=10, cmap='viridis')  
plt.scatter(normals[:,0], normals[:,1], s=5, color='k', label='Normal  
    ↪Observations')  
plt.scatter(anomals[:,0], anomals[:,1], s=5, color='r', label='Anomalies  
    ↪(Contamination = 0.1)')  
plt.title('Anomaly Detection using Elliptic Envelope')  
plt.colorbar(cntr)  
plt.legend()  
plt.grid()  
plt.show()  
  
envelope_anomalies = df.iloc[np.where(y_pred == -1)[0]]
```



1.V.b Kernel Density

```
[10]: # Source Code is From AI221 Github Repository:  

    ↪AI221\Clustering_Anomaly_Detect\anomaly_detect_methods.ipynb  

# Generate the KDE surface as Z  

X_pca_scaled = scaler.fit_transform(X_pca)  

kde = KernelDensity(kernel='gaussian',bandwidth=0.4).fit(X_pca_scaled)

Xp, Yp = np.meshgrid(  

    np.linspace(X_pca_scaled[:, 0].min() * 1.1, X_pca_scaled[:, 0].max() * 1.1,  

    ↪50),  

    np.linspace(X_pca_scaled[:, 1].min() * 1.1, X_pca_scaled[:, 1].max() * 1.1,  

    ↪50))
XY = np.vstack([Xp.ravel(), Yp.ravel()]).T

Zp = np.exp(kde.score_samples(XY))
Zp = Zp.reshape(Xp.shape)

# Establish a confidence level of 95% (or 5% cutoff)
```

```

# for the UCL using the quantile of kde_scores.
scores = kde.score_samples(X_pca_scaled)
threshold = np.quantile(scores,0.05)
print(f"Threshold (KDE) = {np.exp(threshold)}")

# Get the anomalous data points
normals = X_pca_scaled[scores > threshold,:]
anomals = X_pca_scaled[scores <= threshold,:]

cntr = plt.contourf(Xp, Yp, Zp, levels=np.linspace(Zp.min(), Zp.max(), 50),  

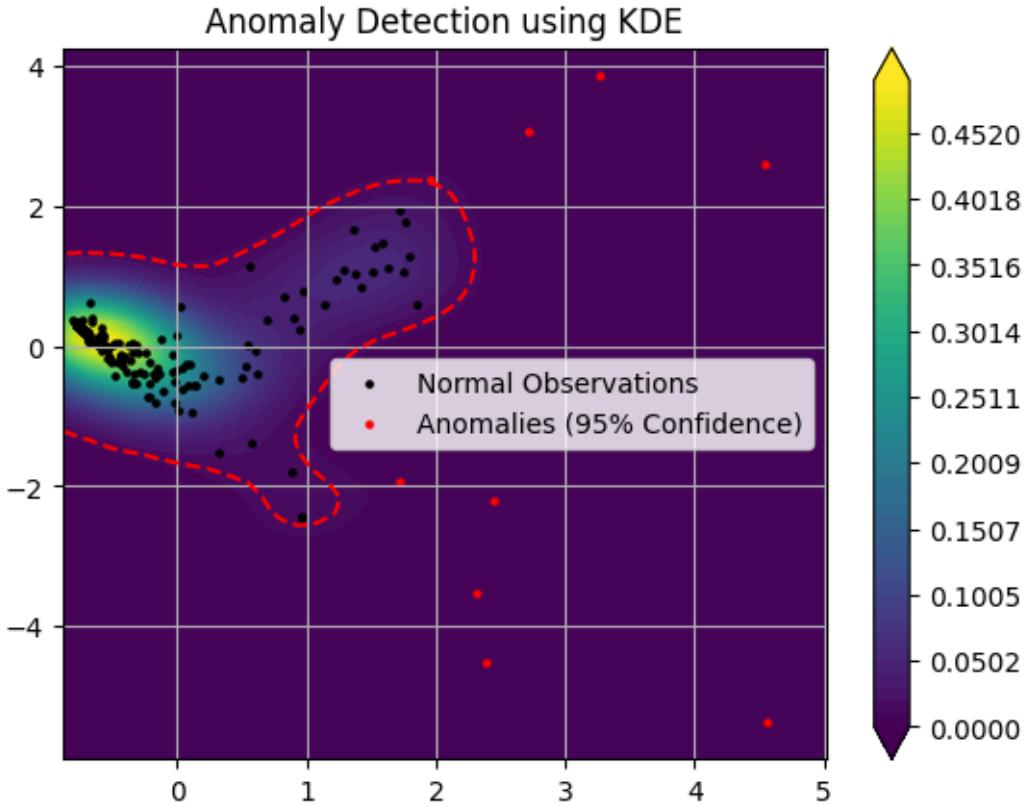
    ↪cmap='viridis', extend='both')
plt.contour(Xp, Yp, Zp, levels=[np.exp(threshold)], colors='red',  

    ↪linestyles='dashed', linewidths=1.5)
plt.scatter(normals[:,0], normals[:,1], s=5, color='k', label='Normal  
Observations')
plt.scatter(anomals[:,0], anomals[:,1], s=5, color='r', label='Anomalies (95%  
Confidence)')
plt.title('Anomaly Detection using KDE')
plt.colorbar(cntr)
plt.legend()
plt.grid()
plt.show()

kernel_density_anomalies = df.iloc[np.where(scores <= threshold)[0]]

```

Threshold (KDE) = 0.012569474629724105



1.V.c. One Class SVM

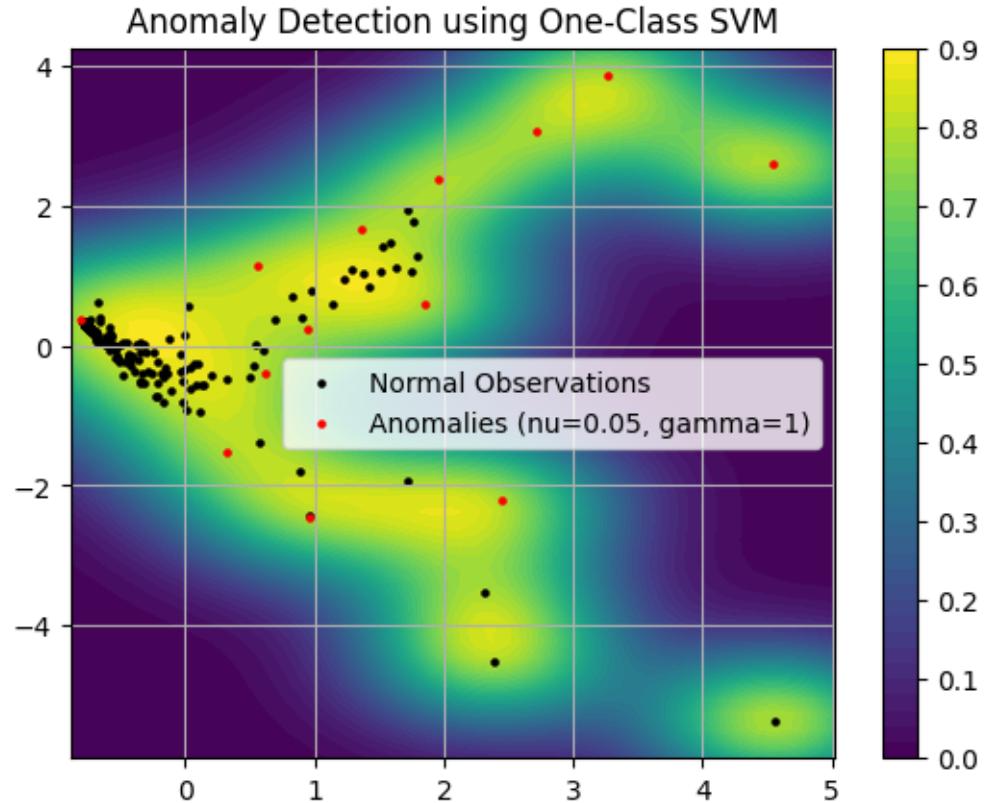
```
[11]: # Source Code is From AI221 Github Repository:  
      ↵AI221\Clustering_Anomaly_Detect\anomaly_detect_methods.ipynb  
ocsvm = OneClassSVM(nu=0.05, gamma=1).fit(X_pca_scaled)  
Zp = ocsvm.score_samples(XY)  
Zp = Zp.reshape(Xp.shape)  
  
# Get the anomalous data points  
y_pred = ocsvm.predict(X_pca_scaled)  
normals = X_pca_scaled[y_pred == 1,:]  
anomals = X_pca_scaled[y_pred == -1,:]  
  
cntr = plt.contourf(Xp, Yp, Zp, levels=50, cmap='viridis')  
plt.scatter(normals[:,0], normals[:,1], s=5, color='k', label='Normal  
↪Observations')  
plt.scatter(anomals[:,0], anomals[:,1], s=5, color='r', label='Anomalies (nu=0.  
↪05, gamma=1)')  
plt.title('Anomaly Detection using One-Class SVM')  
plt.colorbar(cntr)  
plt.legend()
```

```

plt.grid()
plt.show()

ocsvm_anomalies = df.iloc[np.where(y_pred == -1)[0]]

```



1.V.d. Local Outlier Factor

```

[12]: # Source Code is From AI221 Github Repository:
      ↵AI221\Clustering_Anomaly_Detect\anomaly_detect_methods.ipynb
lof = LocalOutlierFactor(n_neighbors=5,novelty=True).fit(X_pca_scaled)
Zp = lof.score_samples(XY)
Zp = Zp.reshape(Xp.shape)

# Get the anomalous data points
y_pred = lof.predict(X_pca_scaled)
normals = X_pca_scaled[y_pred == 1,:]
anomals = X_pca_scaled[y_pred == -1,:]

cntr = plt.contourf(Xp, Yp, Zp, levels=10, cmap='viridis')
plt.scatter(normals[:,0], normals[:,1], s=5, color='k', label='Normal Observations')

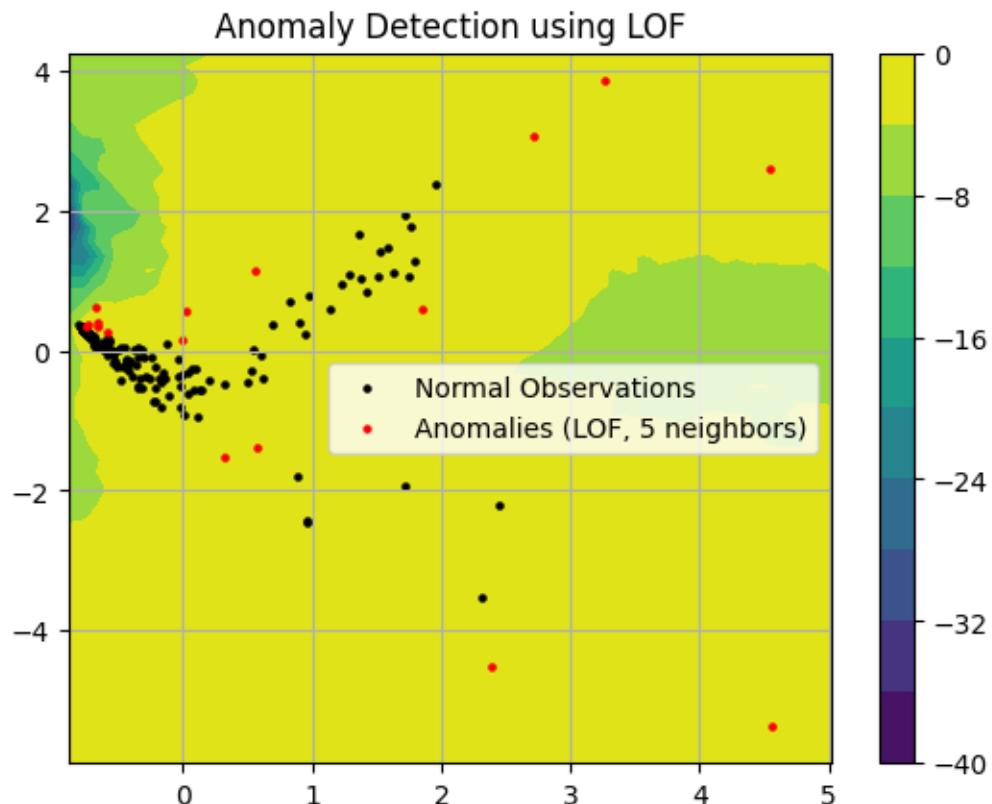
```

```

plt.scatter(anomalies[:,0], anomalies[:,1], s=5, color='r', label='Anomalies (LOF, 5 neighbors)')
plt.title('Anomaly Detection using LOF')
plt.colorbar(cntr)
plt.legend()
plt.grid()
plt.show()

lof_anomalies = df.iloc[np.where(y_pred == -1)[0]]

```



1.V.e. Isolation Forest

```

[13]: # Source Code is From AI221 Github Repository:
      ↪AI221\Clustering_Anomaly_Detect\anomaly_detect_methods.ipynb
isoforest = IsolationForest(contamination=0.1).fit(X_pca_scaled)
Zp = isoforest.score_samples(XY)
Zp = Zp.reshape(Xp.shape)

# Get the anomalous data points
y_pred = isoforest.predict(X_pca_scaled)
normals = X_pca_scaled[y_pred == 1,:]

```

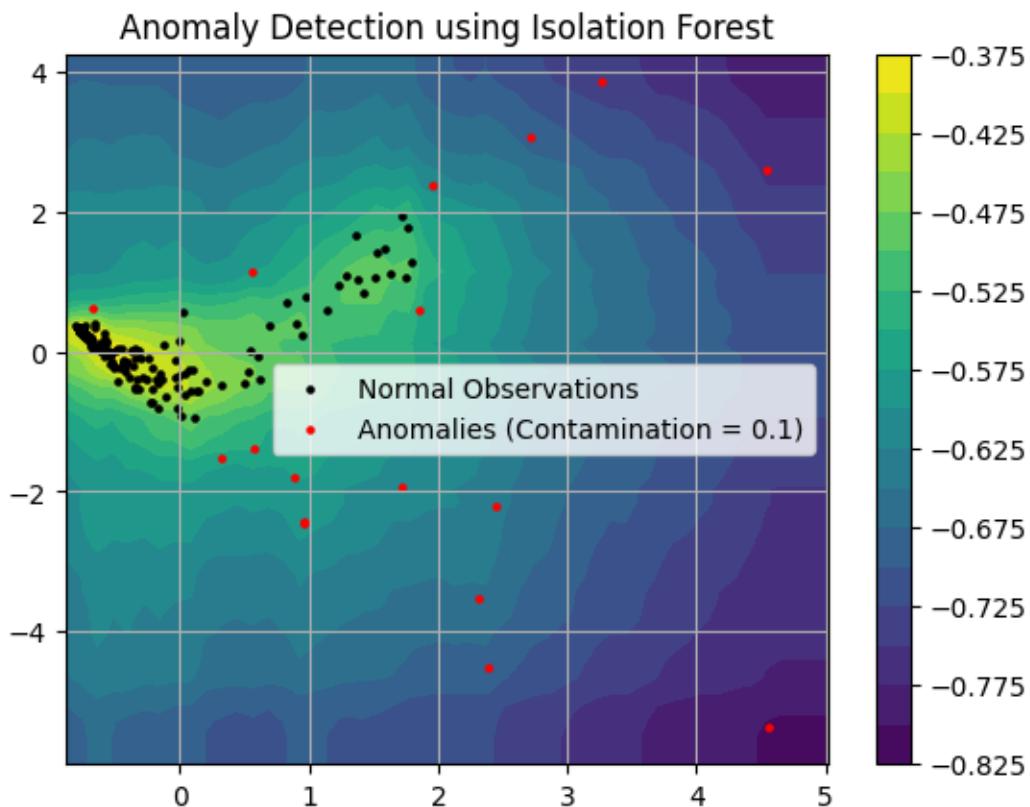
```

anomals = X_pca_scaled[y_pred == -1,:]

cntr = plt.contourf(Xp, Yp, Zp, levels=20, cmap='viridis')
plt.scatter(normals[:,0], normals[:,1], s=5, color='k', label='Normal Observations')
plt.scatter(anomals[:,0], anomals[:,1], s=5, color='r', label='Anomalies (Contamination = 0.1)')
plt.title('Anomaly Detection using Isolation Forest')
plt.colorbar(cntr)
plt.legend()
plt.grid()
plt.show()

isoforest_anomalies = df.iloc[np.where(y_pred == -1)[0]]

```



Item 1.e. Which countries are deemed to be outliers and why? Which features make them outliers?

Based on the collective results from various anomaly detection models, the analysis reveals the following: Switzerland, Luxembourg, and Norway are identified as anomalies by all five models. Denmark is flagged by four models, while Japan, Brunei, Qatar, Singapore, Libya, the United States, and the Bahamas are identified by three models.

A statistical comparison between the combined anomalies from the five models and the overall dataset highlights significant differences. The anomalies exhibit higher values in exports, imports, income, and GDP per capita (GDPP) compared to the overall dataset. This indicates that the countries identified as anomalies are developed nations.

In contrast, the majority of countries in the dataset are either underdeveloped or still developing. Alternatively, it could suggest that the anomaly countries are exceptionally successful, standing out significantly from the rest.

The following code provides supporting evidence for the analysis described above.

```
[14]: # Prompt: Create a bar plot showing the count of anomalies detected per country
# across all methods, x-axis: country, y-axis: count of anomalies
import pandas as pd

import matplotlib.pyplot as plt

# Combine all anomaly dataframes into one
all_anomalies = pd.concat([
    envelope_anomalies['country'],
    kernel_density_anomalies['country'],
    ocsvm_anomalies['country'],
    lof_anomalies['country'],
    isoforest_anomalies['country']
])

# Count the occurrences of each country
anomaly_counts = all_anomalies.value_counts()

# Plot the bar plot
plt.figure(figsize=(12, 8))
anomaly_counts.plot(kind='bar', color='skyblue', edgecolor='black')
plt.title('Anomaly Count per Country', fontsize=16)
plt.xlabel('Country', fontsize=14)
plt.ylabel('Anomaly Count', fontsize=14)
plt.xticks(rotation=90, fontsize=10)
plt.tight_layout()
plt.show()

all_anomalies_df = pd.concat([
    envelope_anomalies,
    kernel_density_anomalies,
    ocsvm_anomalies,
    lof_anomalies,
    isoforest_anomalies
])

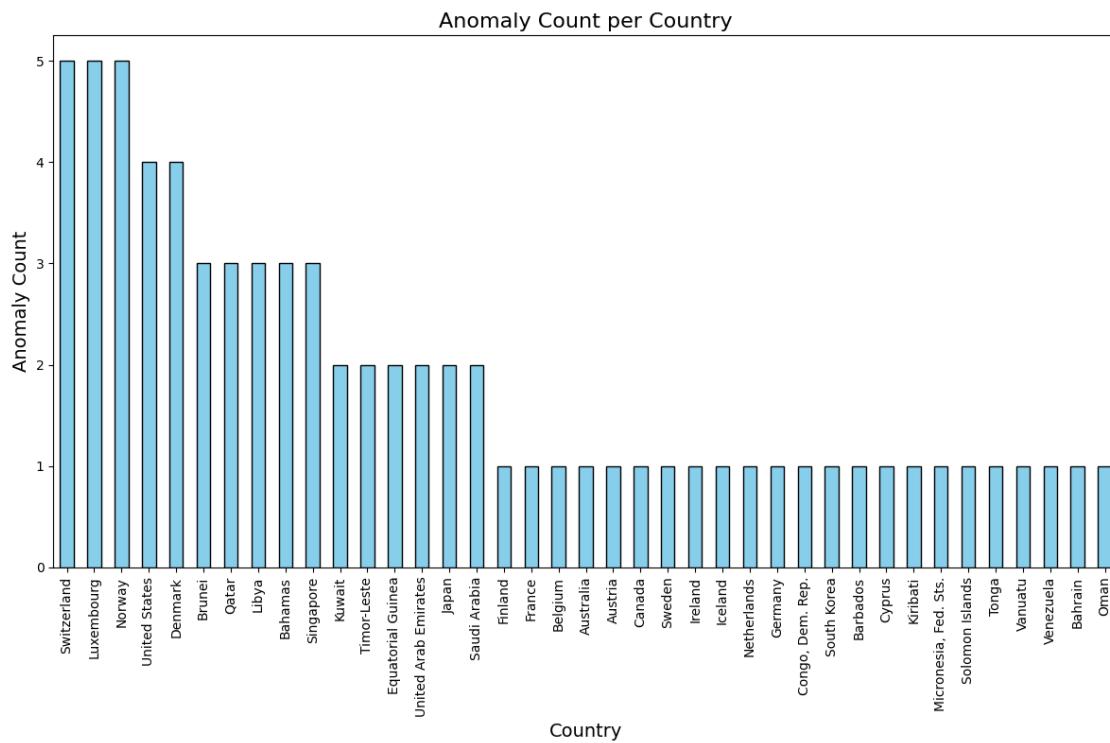
# Prompt: Remove duplicate entries based on country from all_anomalies_df
```

```

all_anomalies_df = all_anomalies_df.drop_duplicates(subset=['country'])
    ↪sort_values(by='country')

print(display(all_anomalies_df))
# Prompt: Show the Mean and Standard Deviation of each feature in
    ↪all_anomalies_df
anomaly_stats = all_anomalies_df.describe().loc[['mean', 'std']]
print(display(anomaly_stats))
# Prompt: Show the Mean and Standard Deviation of each feature in the original
    ↪dataframe df
original_stats = df.describe().loc[['mean', 'std']]
print(display(original_stats))

```



	country	child_mort	exports	health	imports	income	\
7	Australia	4.8	19.8	8.73	20.9	41400	
8	Austria	4.3	51.3	11.00	47.8	43200	
10	Bahamas	13.8	35.0	7.89	43.7	22900	
11	Bahrain	8.6	69.5	4.97	50.9	41100	
13	Barbados	14.2	39.5	7.97	48.7	15300	
15	Belgium	4.5	76.4	10.70	74.7	41100	
23	Brunei	10.5	67.4	2.84	28.0	80600	
29	Canada	5.6	29.1	11.30	31.0	40700	
37	Congo, Dem. Rep.	116.0	41.1	7.91	49.6	609	

42	Cyprus	3.6	50.2	5.97	57.5	33900
44	Denmark	4.1	50.5	11.40	43.6	44000
49	Equatorial Guinea	111.0	85.8	4.48	58.9	33700
53	Finland	3.0	38.7	8.95	37.4	39800
54	France	4.2	26.8	11.90	28.1	36900
58	Germany	4.2	42.3	11.60	37.1	40400
68	Iceland	2.6	53.4	9.40	43.3	38800
73	Ireland	4.2	103.0	9.19	86.5	45700
77	Japan	3.2	15.0	9.49	13.6	35800
81	Kiribati	62.7	13.3	11.30	79.9	1730
82	Kuwait	10.8	66.7	2.63	30.4	75200
89	Libya	16.6	65.6	3.88	42.1	29600
91	Luxembourg	2.8	175.0	7.77	142.0	91700
101	Micronesia, Fed. Sts.	40.0	23.5	14.20	81.0	3340
110	Netherlands	4.5	72.0	11.90	63.6	45500
114	Norway	3.2	39.7	9.48	28.5	62300
115	Oman	11.7	65.7	2.77	41.2	45300
123	Qatar	9.0	62.3	1.81	23.8	125000
128	Saudi Arabia	15.7	49.6	4.29	33.0	45400
133	Singapore	2.8	200.0	3.96	174.0	72100
136	Solomon Islands	28.1	49.3	8.55	81.2	1780
138	South Korea	4.1	49.4	6.93	46.2	30400
144	Sweden	3.0	46.2	9.63	40.7	42900
145	Switzerland	4.5	64.0	11.50	53.3	55500
149	Timor-Leste	62.6	2.2	9.12	27.8	1850
151	Tonga	17.4	12.4	5.07	60.3	4980
157	United Arab Emirates	8.6	77.7	3.66	63.6	57600
159	United States	7.3	12.4	17.90	15.8	49400
162	Vanuatu	29.2	46.6	5.25	52.7	2950
163	Venezuela	17.1	28.5	4.91	17.6	16500

	inflation	life_expec	total_fer	gdpp
7	1.160	82.0	1.93	51900
8	0.873	80.5	1.44	46900
10	-0.393	73.8	1.86	28000
11	7.440	76.0	2.16	20700
13	0.321	76.7	1.78	16000
15	1.880	80.0	1.86	44400
23	16.700	77.1	1.84	35300
29	2.870	81.3	1.63	47400
37	20.800	57.5	6.54	334
42	2.010	79.9	1.42	30800
44	3.220	79.5	1.87	58000
49	24.900	60.9	5.21	17100
53	0.351	80.0	1.87	46200
54	1.050	81.4	2.03	40600
58	0.758	80.1	1.39	41800
68	5.470	82.0	2.20	41900

73	-3.220	80.4	2.05	48700
77	-1.900	82.8	1.39	44500
81	1.520	60.7	3.84	1490
82	11.200	78.2	2.21	38500
89	14.200	76.1	2.41	12100
91	3.620	81.3	1.63	105000
101	3.800	65.4	3.46	2860
110	0.848	80.7	1.79	50300
114	5.950	81.0	1.95	87800
115	15.600	76.1	2.90	19300
123	6.980	79.5	2.07	70300
128	17.200	75.1	2.96	19300
133	-0.046	82.7	1.15	46600
136	6.810	61.7	4.24	1290
138	3.160	80.1	1.23	22100
144	0.991	81.5	1.98	52100
145	0.317	82.2	1.52	74600
149	26.500	71.1	6.23	3600
151	3.680	69.9	3.91	3550
157	12.500	76.5	1.87	35000
159	1.220	78.7	1.93	48400
162	2.620	63.0	3.50	2970
163	45.900	75.4	2.47	13500

None

	child_mort	exports	health	imports	income	inflation	\
mean	17.541026	54.279487	8.005128	51.282051	39408.692308	6.893846	
std	26.812622	38.638883	3.608391	31.583556	26523.865636	9.778181	

	life_expec	total_fer	gdpp
mean	76.123077	2.454359	35158.820513
std	7.033995	1.286196	24695.282429

None

	child_mort	exports	health	imports	income	inflation	\
mean	38.270060	41.108976	6.815689	46.890215	17144.688623	7.781832	
std	40.328931	27.412010	2.746837	24.209589	19278.067698	10.570704	

	life_expec	total_fer	gdpp
mean	70.555689	2.947964	12964.155689
std	8.893172	1.513848	18328.704809

None

1.2 2. Early-Stage Diabetes Risk Prediction

1.2.1 2.I. Load Dataset

```
[19]: # code source: https://archive.ics.uci.edu/dataset/529/
    ↪early+stage+diabetes+risk+prediction+dataset
from ucimlrepo import fetch_ucirepo

# fetch dataset
early_stage_diabetes_risk_prediction = fetch_ucirepo(id=529)

# data (as pandas dataframes)
X = early_stage_diabetes_risk_prediction.data.features
y = early_stage_diabetes_risk_prediction.data.targets

# metadata
print(early_stage_diabetes_risk_prediction.metadata)

# variable information
print(early_stage_diabetes_risk_prediction.variables)
```

{'uci_id': 529, 'name': 'Early Stage Diabetes Risk Prediction',
'repository_url': 'https://archive.ics.uci.edu/dataset/529/early+stage+diabetes+
risk+prediction+dataset', 'data_url':
'https://archive.ics.uci.edu/static/public/529/data.csv', 'abstract': 'This
dataset contains the sign and symptom data of newly diabetic or would be
diabetic patient.', 'area': 'Computer Science', 'tasks': ['Classification'],
'characteristics': ['Multivariate'], 'num_instances': 520, 'num_features': 16,
'feature_types': ['Categorical', 'Integer'], 'demographics': ['Age', 'Gender'],
'target_col': ['class'], 'index_col': None, 'has_missing_values': 'no',
'missing_values_symbol': None, 'year_of_dataset_creation': 2020, 'last_updated':
'Mon Mar 04 2024', 'dataset_doi': '10.24432/C5VG8H', 'creators': [],
'intro_paper': {'ID': 397, 'type': 'NATIVE', 'title': 'Likelihood Prediction of
Diabetes at Early Stage Using Data Mining Techniques', 'authors': 'M. M. F.
Islam, Rahatara Ferdousi, Sadikur Rahman, Humayra Yasmin Bushra', 'venue':
'Computer Vision and Machine Intelligence in Medical Image Analysis', 'year':
2019, 'journal': None, 'DOI': '10.1007/978-981-13-8798-2_12', 'URL': 'https://ww
w.semanticscholar.org/paper/9329dec57c5f13f195220ffa7077fd0029983f07', 'sha':
None, 'corpus': None, 'arxiv': None, 'mag': None, 'acl': None, 'pmid': None,
'pmcid': None}, 'additional_info': {'summary': 'This has been col-\r\nlected
using direct questionnaires from the patients of Sylhet Diabetes\r\nHospital in
Sylhet, Bangladesh and approved by a doctor.'}, 'purpose': None, 'funded_by':
None, 'instances_represent': None, 'recommended_data_splits': None,
'sensitive_data': None, 'preprocessing_description': None, 'variable_info': 'Age
1.20-65\t\t\r\nSex 1. Male, 2.Female\t\t\r\nPolyuria 1.Yes,
2.No.\t\t\r\nPolydipsia 1.Yes, 2.No.\t\t\r\nsudden weight loss 1.Yes,
2.No.\t\t\r\nweakness 1.Yes, 2.No.\t\t\r\nPolyphagia 1.Yes, 2.No.\t\t\r\nGenital
thrush 1.Yes, 2.No.\t\t\r\nvisual blurring 1.Yes, 2.No.\t\t\r\nItching 1.Yes,
2.No.\t\t\r\nIrritability 1.Yes, 2.No.\t\t\r\nDelayed healing 1.Yes,

```

2.No.\t\t\r\npartial paresis 1.Yes, 2.No.\t\t\r\nmuscle sti\x0bness 1.Yes,
2.No.\t\t\r\nAlopecia 1.Yes, 2.No.\t\t\r\nObesity 1.Yes, 2.No.\t\t\r\nClass
1.Positive, 2.Negative.\t\t\r\n', 'citation': None}]

      name      role      type demographic description units \
0        age    Feature    Integer       Age      None  None
1     gender    Feature  Categorical   Gender      None  None
2   polyuria    Feature     Binary     None      None  None
3  polydipsia    Feature     Binary     None      None  None
4 sudden_weight_loss    Feature     Binary     None      None  None
5     weakness    Feature     Binary     None      None  None
6   polyphagia    Feature     Binary     None      None  None
7  genital_thrush    Feature     Binary     None      None  None
8  visual_blurring    Feature     Binary     None      None  None
9      itching    Feature     Binary     None      None  None
10   irritability    Feature     Binary     None      None  None
11  delayed_healing    Feature     Binary     None      None  None
12  partial_paresis    Feature     Binary     None      None  None
13  muscle_stiffness    Feature     Binary     None      None  None
14     alopecia    Feature     Binary     None      None  None
15     obesity    Feature     Binary     None      None  None
16       class    Target     Binary     None      None  None

missing_values
0      no
1      no
2      no
3      no
4      no
5      no
6      no
7      no
8      no
9      no
10     no
11     no
12     no
13     no
14     no
15     no
16     no

```

1.2.2 2.II. Encoding, and Dataset Split

Item 2.a. Make the necessary encoding for categorical inputs. Split the data into 80% Training and 20% Testing with stratification.

```
[33]: # Encode categorical variables
from sklearn.preprocessing import LabelEncoder
```

```

X = pd.get_dummies(X, drop_first=True, dtype=int)
# Encode y

le = LabelEncoder()
y = le.fit_transform(y)

# Split the dataset into training and testing sets (stratify = y)
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,random_state=42, stratify=y)

```

1.2.3 2.III. Best Model Using

Item 2.b. Using Optuna, find the best model between the MLP Classifier, Random Forest Classifier, XGBoost Classifier, Logistic Regression, Naïve Bayes Classifier, SVM Classifier (SVC), and kNN Classifier. Set Optuna to maximize the 10-fold cross-validation score (cross_val_score). You are free to design the search space for hyper-parameters in these models.

```

[37]: import optuna
from sklearn.model_selection import cross_val_score
from sklearn.neural_network import MLPClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, f1_score

# Set the logging level to WARNING to suppress info logs
optuna.logging.set_verbosity(optuna.logging.WARNING)

def objective(trial):
    classifier_name = trial.suggest_categorical("classifier",
                                                ["MLP", "RandomForest", "XGBoost", "LogisticRegression", "NaiveBayes", "SVM", "kNN"])

    if classifier_name == "MLP":
        hidden_layer_sizes_choices = [(50,), (100,), (50, 50)]
        params = {
            "hidden_layer_sizes": eval(trial.suggest_categorical("hidden_layer_sizes", hidden_layer_sizes_choices)),
            "activation": trial.suggest_categorical("activation", ["relu", "tanh"]),
            "solver": trial.suggest_categorical("solver", ["adam", "sgd"]),
        }
    else:
        params = {}

```

```

        "alpha": trial.suggest_float("alpha", 1e-5, 1e-1, log=True),
    }
    model = MLPClassifier(**params, max_iter=1000, random_state=42, ↴
↳early_stopping=True)

    elif classifier_name == "RandomForest":
        params = {
            "n_estimators": trial.suggest_int("n_estimators", 50, 300, step=50),
            "max_depth": trial.suggest_int("max_depth", 3, 20),
            "min_samples_split": trial.suggest_int("min_samples_split", 2, 10),
        }
        model = RandomForestClassifier(**params, random_state=42)

    elif classifier_name == "XGBoost":
        params = {
            "max_depth": trial.suggest_int("max_depth", 3, 15),
            "learning_rate": trial.suggest_float("learning_rate", 0.01, 0.3, ↴
↳log=True),
            "n_estimators": trial.suggest_int("n_estimators", 50, 300, step=50),
        }
        model = XGBClassifier(**params, eval_metric="logloss", random_state=42)

    elif classifier_name == "LogisticRegression":
        params = {
            "C": trial.suggest_float("C", 1e-3, 1e3, log=True),
            "solver": trial.suggest_categorical("k-solver", ["liblinear", ↴
↳"lbfgs"]),
        }
        model = LogisticRegression(C=params["C"], solver=params["solver"], ↴
↳max_iter=1000, random_state=42)

    elif classifier_name == "NaiveBayes":
        model = GaussianNB()

    elif classifier_name == "SVM":
        params = {
            "C": trial.suggest_float("C", 1e-3, 1e3, log=True),
            "kernel": trial.suggest_categorical("kernel", ["linear", "rbf", ↴
↳"poly"]),
        }
        model = SVC(**params, random_state=42)

    elif classifier_name == "kNN":
        params = {
            "n_neighbors": trial.suggest_int("n_neighbors", 3, 15),
            "weights": trial.suggest_categorical("weights", ["uniform", ↴
↳"distance"]),
        }

```

```

        }
    model = KNeighborsClassifier(**params)

    # Perform 10-fold cross-validation
    score = cross_val_score(model, X_train, y_train, cv=10, scoring="accuracy").
    ↪mean()

    return score

# Create and run the Optuna study
study = optuna.create_study(direction="maximize", sampler=optuna.samplers.
    ↪TPESampler(seed=42))
study.optimize(objective, n_trials=500, show_progress_bar=True)

# Get the best model and hyperparameters
best_params = study.best_params
best_classifier = best_params.pop("classifier")

if best_classifier == "MLP":
    hidden_layer_sizes = eval(best_params["hidden_layer_sizes"])
    best_model = MLPClassifier(**best_params, max_iter=1000, random_state=42,
    ↪early_stopping=True)
elif best_classifier == "RandomForest":
    best_model = RandomForestClassifier(**best_params, random_state=42)
elif best_classifier == "XGBoost":
    best_model = XGBClassifier(**best_params, eval_metric="logloss",
    ↪random_state=42)
elif best_classifier == "LogisticRegression":
    best_model = LogisticRegression(C=best_params["C"],
    ↪solver=best_params["k-solver"], random_state=42)
elif best_classifier == "NaiveBayes":
    best_model = GaussianNB()
elif best_classifier == "SVM":
    best_model = SVC(**best_params, random_state=42)
elif best_classifier == "kNN":
    best_model = KNeighborsClassifier(**best_params)

# Train the best model on the training data
best_model.fit(X_train, y_train)

# Evaluate on the test data
y_pred = best_model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred, average="weighted")

print(f"Best Classifier: {best_classifier}")
print(f"Best Hyperparameters: {best_params}")
print(f"Test Accuracy: {accuracy}")

```

```

print(f"Test F1-Score: {f1}")

0%|          | 0/500 [00:00<?, ?it/s]

Best Classifier: RandomForest
Best Hyperparameters: {'n_estimators': 50, 'max_depth': 17, 'min_samples_split': 2}
Test Accuracy: 0.9807692307692307
Test F1-Score: 0.9808511271925906

```

Item 2.b. What is the Accuracy and F1-score on the Test Data of the best model?

The best classifier is RandomForest (`n_estimators=50`, `max_depth=17`, `min_samples_split=2`), havning Test Accuracy of 98% and Test F1-Score of 98%.

1.2.4 2.IV. Random Forest

This section focuses on two key objectives: optimizing the Random Forest model through hyper-parameter tuning and comparing its performance with the paper's reported best Random Forest F1-score of 98%.

Item 2.c. In the paper, the best model was found to be Random Forest, having a weighted average F1 score of 0.98. Using your own hyper-parameter search, can you find a better Random Forest model with higher F1 score?

```

[59]: def objective_rf(trial):

    params = {
        "n_estimators": trial.suggest_int("n_estimators", 10, 1000, step=3),
        "max_depth": trial.suggest_int("max_depth", 3, 200),
        "min_samples_split": trial.suggest_int("min_samples_split", 2, 50),
    }
    model = RandomForestClassifier(**params, random_state=42)

    # Perform 10-fold cross-validation
    score = cross_val_score(model, X_train, y_train, cv=10, scoring="accuracy").
    ↪mean()

    return score

# Create and run the Optuna study
study_rf = optuna.create_study(direction="maximize", sampler=optuna.samplers.
    ↪TPESampler(seed=42))
study_rf.optimize(objective_rf, n_trials=100, show_progress_bar=True)

# Get the best model and hyperparameters
best_params_rf = study_rf.best_params

best_model_rf = RandomForestClassifier(**best_params_rf, random_state=42)

```

```

# Train the best model on the training data
best_model_rf.fit(X_train, y_train)

# Evaluate on the test data
y_pred_rf = best_model_rf.predict(X_test)
accuracy_rf = accuracy_score(y_test, y_pred_rf)
f1_rf = f1_score(y_test, y_pred_rf, average="weighted")

print(f"Best Hyperparameters: {best_params_rf}")
print(f"Test Accuracy: {accuracy_rf}")
print(f"Test F1-Score: {f1_rf}")

```

0%| 0/100 [00:00<?, ?it/s]

Best Hyperparameters: {'n_estimators': 985, 'max_depth': 117, 'min_samples_split': 2}

Test Accuracy: 0.9903846153846154

Test F1-Score: 0.9904061137656939

Item 2.c. Using your own hyper-parameter search, can you find a better Random Forest model with higher F1 score?

Improvement to 99%, the optimized Random Forest model (`n_estimators=985, max_depth=117, min_samples_split=2`) produced above show a slight increase in the F1-score of 98% compared to the result reported in the paper.

Item 2.d. Based on the best Random Forest model in item (c), perform any feature importance method to explain the model. What insights can we get from the model?

The results indicate that the top three features significantly influencing the model's performance are: `polyuria_Yes`, `polydipsia_Yes`, and `gender_Male`. These features play a critical role in the model's predictions and provide valuable insights into the factors driving its performance.

Insights:

- Polyuria (Excessive Urination): This is identified as a key feature, highlighting its significance as an early symptom of diabetes.
- Polydipsia (Excessive Thirst): Similarly, this feature underscores its importance as another early indicator of diabetes.
- Gender (Male): While being male appears as a top feature, this likely reflects the dataset's distribution, where a higher proportion of diabetes cases are observed in males. This does not necessarily imply a causal relationship but rather a pattern in the data.

The following blocks of code support the answer provided:

```
[55]: from sklearn.inspection import permutation_importance

# Calculate permutation importance
```

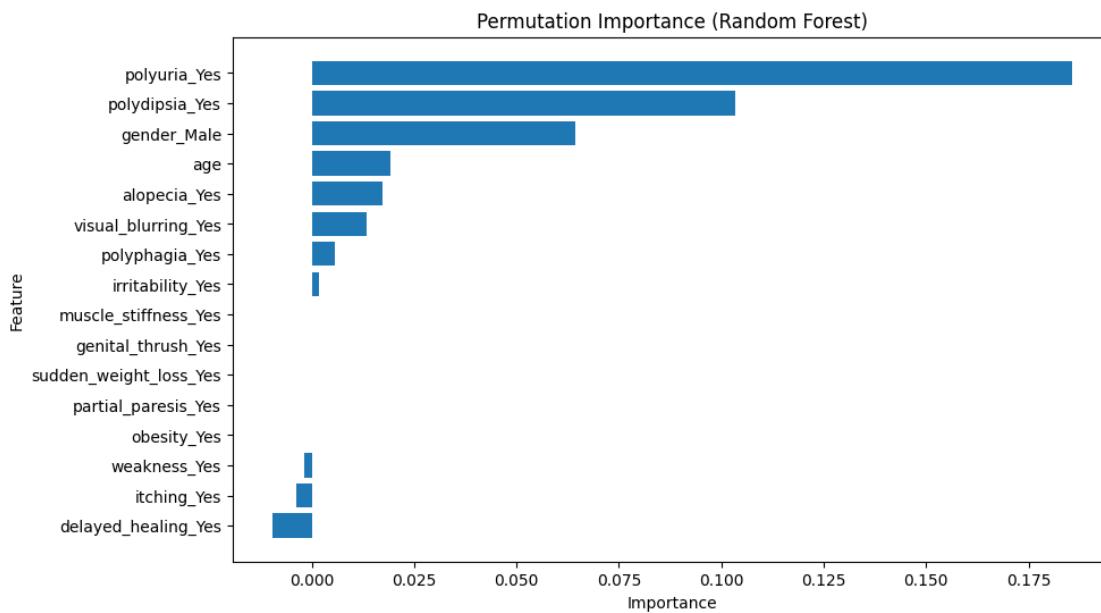
```

perm_importance = permutation_importance(best_model_rf, X_test, y_test, scoring="f1_weighted")

# Create a DataFrame for better visualization
perm_importance_df = pd.DataFrame({
    "Feature": X_train.columns,
    "Importance": perm_importance.importances_mean
}).sort_values(by="Importance", ascending=False)

# Plot the permutation importance
plt.figure(figsize=(10, 6))
plt.barh(perm_importance_df["Feature"], perm_importance_df["Importance"])
plt.xlabel("Importance")
plt.ylabel("Feature")
plt.title("Permutation Importance (Random Forest)")
plt.gca().invert_yaxis()
plt.show()

```



```

[56]: import matplotlib.pyplot as plt
import pandas as pd

# Get feature importance from the trained Random Forest model
importances = best_model_rf.feature_importances_

# Create a DataFrame for better visualization
feature_importance_df = pd.DataFrame({

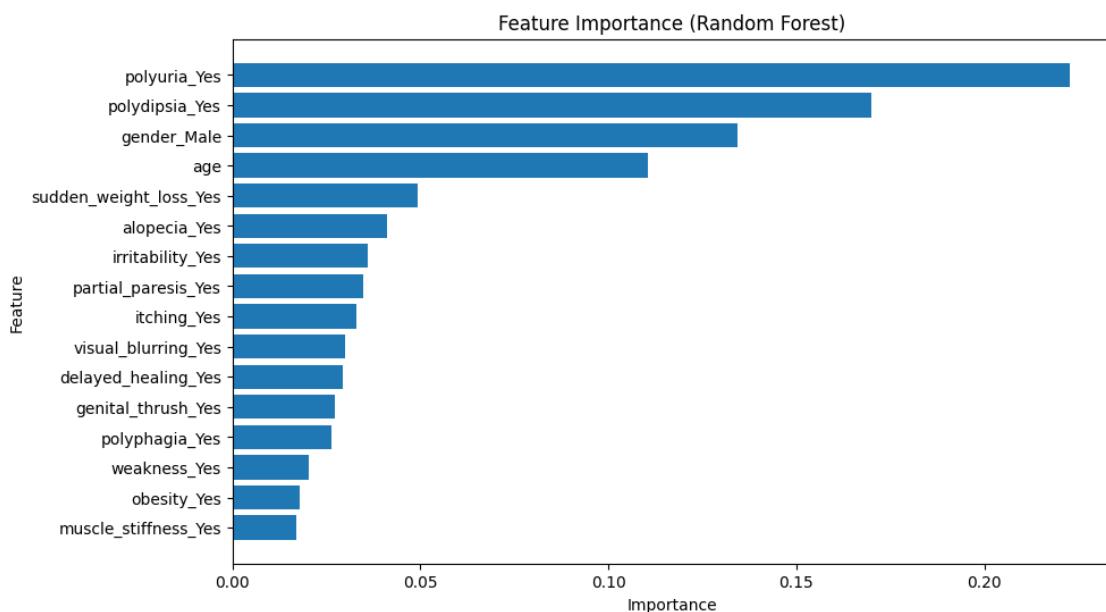
```

```

    "Feature": X_train.columns,
    "Importance": importances
}).sort_values(by="Importance", ascending=False)

# Plot the feature importance
plt.figure(figsize=(10, 6))
plt.barh(feature_importance_df["Feature"], feature_importance_df["Importance"])
plt.xlabel("Importance")
plt.ylabel("Feature")
plt.title("Feature Importance (Random Forest)")
plt.gca().invert_yaxis()
plt.show()

```



```

[57]: # This snippet of code is from MEX1 (household electricity consumption)
from sklearn.pipeline import Pipeline
from sklearn.linear_model import Ridge
from sklearn.metrics import r2_score, accuracy_score

alphas = np.logspace(-3, 3, 10)
best_alpha = None
best_val_score = -np.inf
best_model = None

X_val = X_test
y_val = y_test
for alpha in alphas:

```

```

pipe = Pipeline([
    ('scaler', StandardScaler()),
    ('ridge', Ridge(alpha=alpha))
])
pipe.fit(X_train, y_train)
val_pred = pipe.predict(X_val)
val_score = r2_score(y_val, val_pred)
if val_score > best_val_score:
    best_val_score = val_score
    best_alpha = alpha
    best_model = pipe

# Final evaluation
train_score = r2_score(y_train, best_model.predict(X_train))
test_score = r2_score(y_test, best_model.predict(X_test))

ridge = best_model.named_steps['ridge']

# Get the feature names for X1 to X8
feature_names = np.array(X.columns[:8])

# Get absolute value of coefficients for the best Ridge model
coef_abs = np.abs(ridge.coef_[:8])

# Get indices of top 5 features
top5_idx = np.argsort(coef_abs)[-5:][::-1]
top5_features = feature_names[top5_idx]

print("Top 5 features among X1 to X8:", list(top5_features))

```

Top 5 features among X1 to X8: ['polyuria_Yes', 'polydipsia_Yes', 'gender_Male', 'genital_thrush_Yes', 'sudden_weight_loss_Yes']

1.2.5 2.V. Random forest on larger n_estimators

This section is an additional task where I explore the impact of using a larger value for n_estimators.

```

[ ]: def objective_rf_estimators(trial):
    params = {
        "n_estimators": trial.suggest_int("n_estimators", 5000, 20000, step=500),
        "max_depth": trial.suggest_int("max_depth", 3, 200),
        "min_samples_split": trial.suggest_int("min_samples_split", 2, 50),
    }
    model = RandomForestClassifier(**params, random_state=42)

    # Perform 10-fold cross-validation

```

```

score = cross_val_score(model, X_train, y_train, cv=10, scoring="accuracy").
mean()
return score
# Create and run the Optuna study
study_rf_estimators = optuna.create_study(direction="maximize", sampler=optuna.
samplers.TPESampler(seed=42))
study_rf_estimators.optimize(objective_rf_estimators, n_trials=10, □
show_progress_bar=True)
# Get the best model and hyperparameters
best_params_rf_estimators = study_rf_estimators.best_params

best_model_rf_estimators = RandomForestClassifier(**best_params_rf_estimators, □
random_state=42)
# Train the best model on the training data
best_model_rf_estimators.fit(X_train, y_train)

# Evaluate on the test data
y_pred_rf_estimators = best_model_rf_estimators.predict(X_test)
accuracy_rf_estimators = accuracy_score(y_test, y_pred_rf_estimators)
f1_rf_estimators = f1_score(y_test, y_pred_rf_estimators, average="weighted")

print(f"Best Hyperparameters: {best_params_rf_estimators}")
print(f"Test Accuracy: {accuracy_rf_estimators}")
print(f"Test F1-Score: {f1_rf_estimators}")

```

0% | 0/10 [00:00<?, ?it/s]